

# Algorithm Overview

## Introduction

Insertion Sort is a fundamental comparison-based sorting algorithm, often taught as an introductory method for understanding sorting and algorithmic design. It operates by building a sorted portion of the array one element at a time, inserting each new element into its correct position relative to the already sorted elements.

## Theoretical Background

The algorithm iterates through the input array from left to right. For each element, it compares the current value with elements in the sorted portion to its left, shifting larger elements one position to the right until the correct insertion point is found. The current element is then placed in its correct position. This process is repeated for all elements, resulting in a fully sorted array.

Insertion Sort is classified as an in-place, stable sorting algorithm. It does not require additional memory proportional to the input size, and it preserves the relative order of equal elements. The algorithm is particularly efficient for small datasets and nearly-sorted arrays, where it can approach linear time performance.

## Algorithm Steps

1. Start with the second element of the array (index 1).
2. Compare the current element with the elements before it.
3. Shift all larger elements one position to the right.
4. Insert the current element into its correct position.
5. Repeat steps 2–4 for all elements in the array.

## Use Cases and Practical Relevance

Insertion Sort is rarely used for large datasets due to its quadratic time complexity in the average and worst cases. However, it is highly effective for small arrays and nearly-sorted data.

## Optimizations for Nearly-Sorted Data

Insertion Sort's inner loop naturally terminates as soon as the correct position is found, making it efficient for nearly-sorted data. This is a built-in property of the algorithm, not a custom optimization in this implementation.

# Complexity Analysis

## Time Complexity of Insertion Sort

Insertion Sort's time complexity depends heavily on the initial order of the input array.

### Best Case (Already Sorted Array)

- Process: For each element, the inner loop checks if the previous element is greater. Since the array is sorted, this check fails immediately, and no shifting occurs.
- Comparisons: One comparison per element (except the first), so  $(n-1)$  comparisons.
- Shifts: Zero shifts, as every element is already in place.
- Total Operations: Linear in  $n$ .
- Big-O Notation:
- Best Case:  $\Omega(n)$ ,  $\Theta(n)$ ,  $O(n)$

### Average Case (Random Array)

- Process: Each new element is compared with about half of the sorted portion on average before finding its place.
- Comparisons: On average, the inner loop runs about  $i/2$  times for each  $i$  from 1 to  $n-1$ .
- Total Comparisons:
- Sum from  $i=1$  to  $n-1$  of  $(i/2) \approx n^2/4$
- Shifts: Similar to comparisons, as each comparison may result in a shift.
- Big-O Notation:
- Average Case:  $\Theta(n^2)$

### Worst Case (Reverse Sorted Array)

- Process: Every new element is smaller than all previous elements, so the inner loop runs  $i$  times for each  $i$ .
- Comparisons:
- Sum from  $i=1$  to  $n-1$  of  $i = n(n-1)/2 \approx n^2/2$
- Shifts: Each element is shifted all the way to the front.
- Big-O Notation:
- Worst Case:  $O(n^2)$

## Summary Table

Case	Comparisons	Shifts	Time Complexity
Best	$n-1$	0	$\Omega(n)$ , $\Theta(n)$ , $O(n)$
Average	$\sim n^2/4$	$\sim n^2/4$	$\Theta(n^2)$
Worst	$\sim n^2/2$	$\sim n^2/2$	$O(n^2)$

## Space Complexity of Insertion Sort

- Auxiliary Space: Insertion Sort is an in-place algorithm. It only uses a constant amount of extra space (for variables like temp and loop counters).
- No Recursion: The algorithm is fully iterative, so there is no stack overhead.
- Big-O Notation:
- Space Complexity:  $O(1)$  for all cases.

## Comparison with Selection Sort

Algorithm	Best Case	Average Case	Worst Case	Space Complexity
Insertion Sort	$\Theta(n)$	$\Theta(n^2)$	$O(n^2)$	$O(1)$
Selection Sort	$\Theta(n^2)$	$\Theta(n^2)$	$O(n^2)$	$O(1)$

Selection Sort always performs  $n(n-1)/2$  comparisons, regardless of input order, and at most  $n$  swaps.

Insertion Sort is much faster on nearly-sorted or small arrays, thanks to its linear best case and stability.

## Recurrence Relations

In the worst case, each element (from the second to the last) must be compared with every element before it to find its correct position. For the  $i$ -th element, this means up to  $i$  comparisons.

This leads to the following recurrence:

- $T(1) = 0$
- $T(n) = T(n-1) + (n-1)$
- $T(n) = (n-1) + (n-2) + \dots + 1 + 0 = n(n-1)/2$

Simplified:

$T(n)$  is proportional to  $n^2$ , so the time complexity is  $O(n^2)$  in the worst case.

# Code Review

## Identification of Inefficient Code Sections

### General Structure:

The Insertion Sort implementation is straightforward and follows the standard algorithm. The main sorting logic is contained in a single method, which iterates through the array and shifts elements as needed to insert each value into its correct position.

### Metrics Tracking:

The code integrates a PerformanceTracker to count comparisons, moves, and array accesses. This is a good practice for empirical analysis, but it does add some overhead to each operation. However, this overhead is necessary for benchmarking and does not affect the algorithm's asymptotic complexity.

### Potential Inefficiencies:

- No Early Exit for Sorted Subarrays:
  - The current implementation does not check if the array is already sorted before starting the main loop. For nearly-sorted or already sorted arrays, the algorithm is naturally efficient, but a pre-check could avoid unnecessary work in some cases (though this would add  $O(n)$  overhead and is rarely used in practice).
- No Use of Binary Search for Insertion Point:
  - The inner loop finds the insertion point by linear search. Using binary search could reduce the number of comparisons from  $O(n^2)$  to  $O(n \log n)$ , but the number of shifts (moves) would remain  $O(n^2)$ . This is a classic trade-off: binary search improves comparisons but not overall time complexity.

## Specific Optimization Suggestions

- Binary Search for Insertion Point:

### Rationale:

Replace the linear search in the inner loop with a binary search to find the correct insertion index. This reduces the number of comparisons, especially for large arrays. Use `Arrays.binarySearch()` or a custom binary search to locate the insertion point, then shift elements as usual. Reduces comparisons to  $O(n \log n)$ , but total time complexity remains  $O(n^2)$  due to shifting.

- Early Exit for Sorted Arrays:

**Rationale:**

Before starting the main loop, check if the array is already sorted. If so, skip sorting entirely. Add a single pass to check if each element is greater than or equal to the previous one. Adds  $O(n)$  overhead, but can save time if sorted arrays are common in your use case.

- Reduce Metrics Overhead (for Production):

**Rationale:**

The metrics tracking is useful for analysis but should be disabled or removed in production code to avoid unnecessary overhead. Use conditional compilation or a debug flag to enable/disable metrics. Cleaner, faster code in production environments.

## **Proposed Improvements for Time and Space Complexity**

- **Time Complexity:**
  - o Using binary search for the insertion point can reduce the number of comparisons, but the overall time complexity remains  $O(n^2)$  due to the need to shift elements.
  - o For large arrays, switching to a more efficient algorithm (hybrid approach) can reduce time complexity to  $O(n \log n)$ .
- **Space Complexity:**
  - o The algorithm is already in-place and uses only  $O(1)$  extra space. No further improvements are needed for space efficiency.

**Conclusion:**

The current Insertion Sort implementation is efficient for small and nearly-sorted arrays, and the code is clean and easy to follow. The main opportunities for optimization are reducing comparisons with binary search, adding an early exit for sorted arrays, and using a hybrid approach for large inputs. Space usage is already optimal.

# Empirical Results

## Performance Table (Time vs Input Size)

The following table summarizes the average runtime (ms/op) of Insertion Sort on different input types and sizes, as measured by JMH benchmarks:

Input Type	100	1,000	10,000	100,000
Nearly Sorted	≈0.001	0.008	0.509	44.561
Random	0.002	0.178	14.840	1789.367
Reverse	0.004	0.355	36.643	10983.204
Sorted	≈0.0001	0.001	0.009	0.080

Units: All times are in milliseconds per operation (ms/op).

Source: Data taken from InsertionSortJmhResults.txt.

## Validation of Theoretical Complexity

- **Best Case (Sorted/Nearly Sorted):**

The measured times for sorted and nearly sorted arrays grow slowly with input size, confirming the expected linear time complexity ( $O(n)$ ) in the best case. For example, sorting 100,000 already sorted elements takes only 0.08 ms, while nearly sorted data takes 44.56 ms—still much faster than random or reverse.

- **Average and Worst Case (Random/Reverse):**

For random and reverse-sorted arrays, the runtime increases much more rapidly as  $n$  grows, matching the expected quadratic time complexity ( $O(n^2)$ ). Sorting 100,000 random elements takes 1,789 ms, and reverse-sorted input takes over 10,983 ms, showing the dramatic impact of input order.

- **Growth Pattern:**

The data clearly shows that as the input size increases by a factor of 10, the runtime for random and reverse inputs increases by roughly 100x, which is characteristic of quadratic growth.

## Analysis of Constant Factors and Practical Performance

- **Constant Factors:**

The actual runtime for small arrays is extremely low (fractions of a millisecond), making Insertion Sort practical for small or nearly sorted datasets. The difference between sorted and nearly sorted inputs is small for small  $n$ , but grows with larger  $n$  due to the extra shifts required.

- **Practical Implications:**

- o For small arrays ( $n \leq 1,000$ ), Insertion Sort is extremely fast regardless of input order.
- o For large arrays, performance is highly sensitive to input order. Sorted and nearly sorted data remain efficient, but random and reverse-sorted data become impractically slow.
- o The algorithm's in-place nature means it uses minimal memory, and the measured garbage collection rates (not shown in the main table) are low, confirming efficient memory usage.

### Memory and GC Metrics for Random Input

Size	Alloc Rate (MB/sec)	GC Count	GC Time (ms)
100	204.55	16	11
1,000	18.95	1	1
10,000	1.80	$\approx 0$	—
100,000	0.21	$\approx 0$	—

### Observations:

- Memory allocation per operation grows linearly with input size, as expected for array-based algorithms.
- GC activity is higher for small arrays (due to frequent short-lived allocations) but becomes negligible for large arrays, indicating efficient memory usage and minimal GC overhead at scale.
- The algorithm does not create significant garbage, confirming its in-place nature and low memory footprint.

### Summary:

The empirical results strongly validate the theoretical analysis. Insertion Sort is ideal for small or nearly sorted arrays, but not suitable for large, unsorted datasets due to its quadratic time complexity.

# Conclusion

This review has evaluated the Insertion Sort implementation from both theoretical and practical perspectives. The code is well-organized, readable, and includes performance tracking, which is helpful for empirical analysis. The implementation correctly handles all standard and edge cases, such as empty arrays and already sorted input.

## Summary of Findings

- **Correctness:** The algorithm produces correct results for all tested input types.
- **Performance:**
  - Very efficient on small or nearly sorted arrays, matching the expected linear best-case time complexity.
  - Runtime increases rapidly for large, random, or reverse-sorted arrays, consistent with quadratic worst-case complexity.
- **Empirical Validation:** Benchmark results confirm the theoretical analysis, showing clear differences in runtime depending on input order.
- **Memory Usage:** The algorithm sorts in place, requiring minimal extra memory. Garbage collection and allocation rates are low.

## Suggested Optimizations

- **Binary Search for Insertion Point**

Use binary search to find where to insert each element. This can reduce the number of comparisons, especially for larger arrays, though the overall time complexity remains quadratic.

- **Early Exit for Sorted Input**

Add a check at the start to detect if the array is already sorted. This can save time when sorted data is common.

- **Optional Performance Tracking**

Allow performance tracking to be enabled or disabled as needed.

## Overall Assessment

The current implementation is robust and reliable for its intended use cases. With the suggested improvements, the algorithm can become even more flexible and efficient for a wider range of scenarios.